
CS 189: INTRODUCTION TO
MACHINE LEARNING

Fall 2017



HOMEWORK 13



Solutions by

JINHONG DU

3033483677

Question 1

(a)

Jinhong Du
jaydu@berkeley.edu

(b)

I certify that all solutions are entirely in my words and that I have not looked at another student's solutions. I have credited all external sources in this write up.

Jinhong Du

Question 2

(a)

\therefore

$$\begin{aligned}\sigma(t) &= \frac{1}{1 + e^{-t}} \\ \lim_{t \rightarrow +\infty} \sigma(t) &= \frac{1}{1 + \lim_{t \rightarrow +\infty} e^{-t}} \\ &= 1 \\ \lim_{t \rightarrow -\infty} \sigma(t) &= \frac{1}{1 + \lim_{t \rightarrow -\infty} e^{-t}} \\ &= 0\end{aligned}$$

and $\sigma(t)$ is bounded by 1

$\therefore \sigma(t)$ is a thresholding function

\therefore

$$f(t) = \text{Relu}(t) - \text{Relu}(t - 1)$$

$$= \begin{cases} 1 & , t \geq 1 \\ t & , 0 < t < 1 \\ 0 & , t \leq 0 \end{cases}$$



$$\lim_{t \rightarrow +\infty} f(t) = 1$$

$$\lim_{t \rightarrow -\infty} f(t) = 0$$


and $f(t)$ is bounded by 1

$\therefore \text{Relu}(t) - \text{Relu}(t - 1)$ is a thresholding function

(b)



Solution (cont.)



(c)

Let

$$\tau_n(t) = \tau(nt)$$

then $\forall t_0 > 0$,

$$\begin{aligned} \lim_{n \rightarrow +\infty} \tau_n(t_0) &= \lim_{n \rightarrow +\infty} \tau(nt_0) \\ &= 1 \end{aligned}$$

$\forall t_0 < 0$,

$$\begin{aligned} \lim_{n \rightarrow +\infty} \tau_n(t_0) &= \lim_{n \rightarrow +\infty} \tau(nt_0) \\ &= 0 \end{aligned}$$

\therefore

$$\lim_{n \rightarrow \infty} \tau_n(t) = \begin{cases} 1 & , t > 0 \\ \tau(0) & , t = 0 \\ 0 & , t < 0 \end{cases}$$

For $t_0 = 0$, \therefore

$$\lim_{t \rightarrow \infty} \tau(nt + n) = 1$$

\therefore

$$S(t) \subseteq cl(\{\tau(nt) \text{ for all } n \in \mathbb{N}\})$$

i.e. threshold functions with appropriately scaled arguments are step functions in the limit.

$\therefore \forall (w', b') \in \mathbb{R}^d \times \mathbb{R}, x \in \mathbb{R}^d$,

$$S(< w', x > + b') = \begin{cases} 1 & , < w', x > + b' \geq 0 \\ 0 & , < w', x > + b' < 0 \end{cases}$$

Solution (cont.)

is a step functions

\therefore

$$\begin{aligned} S(< w', x > + b') &\in cl(\{\tau[n(< w', x > + b')]\text{ for all } n \in \mathbb{N}\}) \\ &\subseteq cl(\{\tau(< w, x > + b)\text{ for some } w, b\}) \end{aligned}$$

(d)

$\therefore c(y')$ is a continuous function on $[-1, 1]$

$\therefore c(y')$ is uniformly continuous on $[-1, 1]$, i.e., $\forall \delta > 0, \exists \epsilon > 0$, s.t. $\forall x_1, x_2 \in [-1, 1], |x_1 - x_2| < \epsilon$,

$$|c(x_1) - c(x_2)| < \delta$$

$\therefore \forall a, b \in \mathbb{R}, a < b$,

$$\mathbf{1}_{[a, b]}(t) = S(t - a) - S(t - b)$$

$\therefore \exists n \in \mathbb{N}^+$, s.t. $\Delta: -1 = x_0 < x_1 < \dots < x_n = 1, x_n - x_{n-1} < \epsilon$ and $\forall i \in \mathbb{N}^+, i \leq n, y' \in [x_{i-1}, x_i]$,

$$\begin{aligned} g(y') &= \min_{t \in [x_{i-1}, x_i]} \{c(t)\} \cdot \mathbf{1}_{[x_{i-1}, x_i]}(y') \\ &= \min_{t \in [x_{i-1}, x_i]} \{c(t)\} \cdot [S(y' - x_{i-1}) - S(y' - x_i)] \end{aligned}$$

$$|g(y') - c(y')| < \delta$$

Therefore, $\forall y' \in [-1, 1]$,

$$|g(y') - c(y')| < \delta$$

Let

$$c(y_i) = \min_{t \in [x_{i-1}, x_i]} \{c(t)\}$$

then

$$\begin{aligned} g(y') &= \sum_{i=1}^n c(y_i) [S(y' - x_{i-1}) - S(y' - x_i)] \\ &= \sum_{\substack{i=1 \\ y_{i-1} < y_i}}^n [c(y_i) - c(y_{i-1})] S(y' - x_{i-1}) \end{aligned}$$

and $P_\delta = \{y_1, \dots, y_n\}$

(e)

$$\begin{aligned}
\sum_i |c(z_i) - c(z_{i-1})| &\leq \lim_{\delta \rightarrow 0} \sum_{\substack{i=1 \\ y_{i-1} < y_i}}^n \frac{|\cos(\|w\|_1 y_i) - \cos(\|w\|_1 y_{i-1})|}{(y_i - y_{i-1})} (y_i - y_{i-1}) \\
&\leq \int_{-1}^1 |c(x)|' dx \\
&= \|w\|_1 \int_{-1}^1 |\sin(\|w\|_1 x)| dx \\
&\leq \|w\|_1 \int_{-1}^1 1 dx \\
&\leq 2\|w\|_1
\end{aligned}$$

where $\delta = \max\{x_i - x_{i-1} : 0 < i \leq n\}$

Therefore, for every $w \neq 0$, we have approximated $\cos(\langle w, x \rangle)$ by a linear combination of step functions having sum of absolute coefficients at most $2\|w\|_1$.

(f)

From (c), we have

$$S(\langle w', x \rangle + b') \in cl(\{\tau(\langle w, x \rangle + b) \text{ for some } w, b\})$$

and from (d), we have $\forall c(y')$ can be approximated by

$$g(y') = \sum_{\substack{i=1 \\ y_{i-1} < y_i}}^n [c(y_i) - c(y_{i-1})] S(y' - x_{i-1})$$

Thus, $\frac{c(y')}{2\|w\|_1}$ can be approximated by

$$g(y') = \sum_{\substack{i=1 \\ y_{i-1} < y_i}}^n \frac{c(y_i) - c(y_{i-1})}{2\|w\|_1} S(y' - x_{i-1})$$

and

$$\begin{aligned}
\left| \sum_{\substack{i=1 \\ y_{i-1} < y_i}}^n \frac{c(y_i) - c(y_{i-1})}{2\|w\|_1} \right| &\leq \frac{1}{2\|w\|_1} \sum_{\substack{i=1 \\ y_{i-1} < y_i}}^n |c(y_i) - c(y_{i-1})| \\
&\leq 1
\end{aligned}$$

i.e., it is bounded.

I.e., $\forall f \in \mathcal{F}$, f is linear combination of step functions and the sum of coefficients equals 1.

Therefore,

$$\mathcal{F} \subseteq \overline{conv}\{\tau(\langle w, x \rangle + b) \text{ for some } w, b\}$$

(g)

\therefore

$$\begin{aligned}
\mathbb{E}G &= \sum_{i=1}^m c_i g_i^* \\
&= f \\
\text{Var}(G) &= \mathbb{E}G^2 - (\mathbb{E}G)^2 \\
&\leq \mathbb{E}G^2 \\
&= \sum_{i=1}^m c_i g_i^{*2} \\
&\leq \sum_{i=1}^m c_i \\
&= 1
\end{aligned}$$

\therefore

$$\begin{aligned}
\mathbb{E} \left[\frac{1}{p} \sum_{i=1}^p G_i \right] &= \frac{1}{p} \sum_{i=1}^p \mathbb{E}(G_i) \\
&= f \\
\text{Var} \left[\frac{1}{p} \sum_{i=1}^p G_i \right] &= \frac{1}{p^2} \sum_{i=1}^p \text{Var} G_i \\
&\leq \frac{1}{p}
\end{aligned}$$

\therefore

$$\begin{aligned}
\mathbb{E}[\|f_p - f\|^2] &= \mathbb{E} \left\{ \int_{x \in [0,1]^d} \left[\frac{1}{p} \sum_{i=1}^p G_i - f(x) \right]^2 dx \right\} \\
&= \int_{x \in [0,1]^d} \mathbb{E} \left[\frac{1}{p} \sum_{i=1}^p G_i - f(x) \right]^2 dx \\
&= \int_{x \in [0,1]^d} \text{Var} \left[\frac{1}{p} \sum_{i=1}^p G_i \right] dx \\
&\leq \int_{x \in [0,1]^d} \frac{1}{p} dx \\
&= \frac{1}{p}
\end{aligned}$$

(h)

\therefore from (f) we have

$$\mathcal{F}_{\cos} \subseteq \overline{\text{conv}}(\{\tau(< w, x > + b) \text{ for some } w, b\})$$

$\therefore \quad \forall \epsilon > 0, \exists c_1, \dots, c_p \in \mathbb{R}, g_1^*, \dots, g_p^*$ where $c_i \geq 0, \sum_{i=1}^p c_i = 1$ and $g_i^* = \tau(< w_i, x > + b_i)$ for some w_i, b_i , s.t. $g = \sum_{i=1}^p c_i g_i^*, |f - g| < \epsilon$

Solution (cont.)

\therefore

$$\begin{aligned}\mathbb{E}[\|f_p - f\|^2] &\leq \mathbb{E}[\|f_p - g\|^2] + \mathbb{E}[\|\epsilon\|^2] \\ &\leq \frac{1}{p} + \epsilon\end{aligned}$$

\therefore


$$\begin{aligned}E(f, p) &= \inf_{h \in \mathcal{H}_p} \int_{x \in [0, 1]^d} [f(x) - h(x)]^2 dx \\ &\leq \mathbb{E}[\|f_p - f\|^2] + \epsilon \\ &\leq \frac{1}{p} + \epsilon\end{aligned}$$

Therefore,

$$E(f, p) \leq \frac{1}{p}$$

Question 3

(a)



```

1  def build_network(self,
2      images,
3      num_outputs,
4      scope='yolo'):
5
6      with tf.variable_scope(scope):
7          with slim.arg_scope([slim.conv2d, slim.fully_connected],
8              weights_initializer=
9                  tf.truncated_normal_initializer(0.0, 0.01),
10                 weights_regularizer=
11                     slim.l2_regularizer(0.0005)):
12
13             self.conv1 = slim.conv2d(images, 5, [15, 15],
14                 activation_fn = None, scope='conv1')
15             relu1 = tf.nn.relu(self.conv1)
16             pool1 = slim.max_pool2d(relu1, [3, 3], scope='pool1')
17             fc2 = slim.fully_connected(slim.flatten(pool1),
18                 512, activation_fn = None, scope='fc2')
19             relu2 = tf.nn.relu(fc2)
20             net = slim.fully_connected(relu2, 25,
21                 activation_fn = None, scope='fc3')
22
23     return net

```

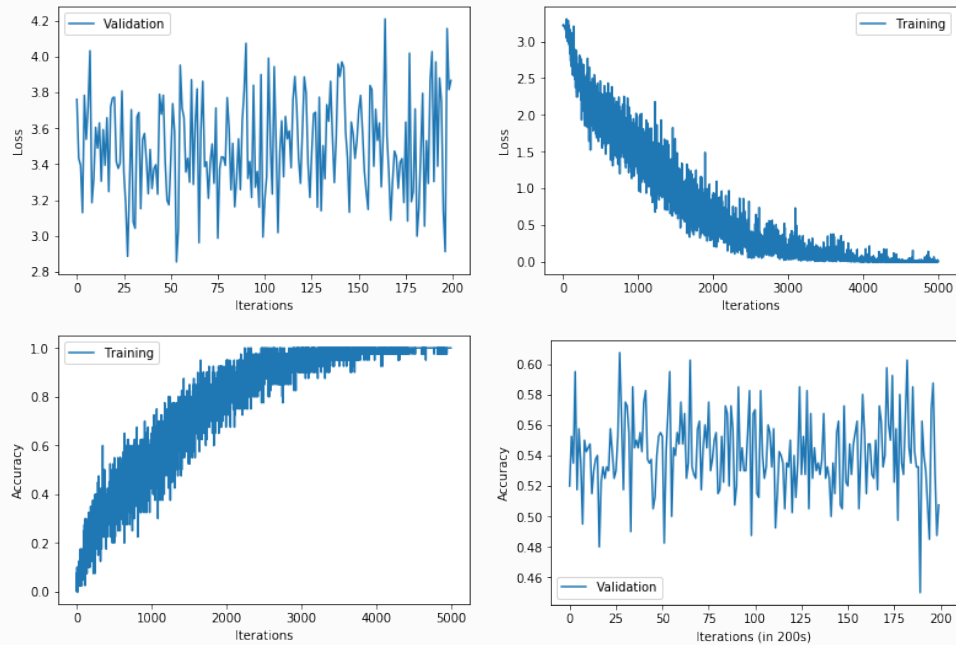
(b)

```


1 def get_train_batch(self):
2     batch = np.random.choice(self.train_data, self.batch_size)
3     features = np.array([i['features'] for i in batch])
4     labels = np.array([i['label'] for i in batch])
5     return features, labels
6
7 def get_validation_batch(self):
8     batch = np.random.choice(self.val_data, self.val_batch_size)
9     features = np.array([i['features'] for i in batch])
10    labels = np.array([i['label'] for i in batch])
11    return features, labels

```


(c)



(d)




Solution (cont.)



```
1 def visualize_features(self, net):
2
3     images = [0, 10, 100]
4     for j in images:
5         batch_eval = np.zeros(
6             [1, self.train_data[0]['features'].shape[0],
7              self.train_data[0]['features'].shape[1],
8              self.train_data[0]['features'].shape[2]])
9         batch_eval[0, :, :, :] = self.train_data[j]['features']
10
11        response_map = self.sess.run(net.conv1,
12                                       feed_dict={net.images: batch_eval})
13
14        for i in range(5):
15            plt.subplot(1, 5, i+1)
16            img = self.revert_image(response_map[0, :, :, i])
17            plt.imshow(img)
18        plt.show()
```

(e)



```
1 class NN():
2     def __init__(self, train_data, val_data, n_neighbors=5):
3         self.train_data = train_data
4         self.val_data = val_data
```

Solution (cont.)

```
5
6     self.sample_size = 400
7     self.model = KNeighborsClassifier(n_neighbors=n_neighbors)
8
9     def train_model(self):
10         x = np.array([np.reshape(self.train_data[i]['features'],
11                                (90*90*3)) for i in range(len(self.train_data))])
12         y = np.array([self.train_data[i]['label'] for i in
13                        range(len(self.train_data))])
14         self.model.fit(x,y)
15
16     def get_validation_error(self):
17         index = np.random.choice(len(self.val_data), self.sample_size,
18                                  replace=False)
19         predy = self.model.predict([np.reshape(
20             self.val_data[i]['features'],
21             (90*90*3)) for i in index])
22         return np.mean(np.argmax(predy,1)!=np.argmax([
23             self.val_data[i]['label'] for i in index],1))
24
25     def get_train_error(self):
26         index = np.random.choice(len(self.train_data),
27                                  self.sample_size, replace=False)
28         predy = self.model.predict([np.reshape(
29             self.train_data[i]['features'],
30             (90*90*3)) for i in index])
31         return np.mean(np.max(predy,1)!=np.argmax([
32             self.train_data[i]['label'] for i in index],1))
```

Question 4

Question What are the regularization methods in CNN?

Solution Empirical

(1) Dropout

Because a fully connected layer occupies most of the parameters, it is prone to overfitting. One method to reduce overfitting is dropout. At each training stage, individual nodes are either “dropped out” of the net with probability $1 - p$ or kept with probability p , so that a reduced network is left; incoming and outgoing edges to a dropped-out node are also removed. Only the reduced network is trained on the data in that stage. The removed nodes are then reinserted into the network with their original weights.

In the training stages, the probability that a hidden node will be dropped is usually 0.5; for input nodes, this should be much lower, intuitively because information is directly lost when input nodes are ignored.

At testing time after training has finished, we would ideally like to find a sample average of all possible 2^n dropped-out networks; unfortunately this is unfeasible for large values of n . However, we can find an approximation by using the full network with each node’s output weighted by a factor of p , so the expected value of the output of any node is the same as in the training stages. This is the biggest contribution of the dropout method: although it effectively generates 2^n neural nets, and as such allows for model combination, at test time only a single network needs to be tested.

By avoiding training all nodes on all training data, dropout decreases overfitting in neural nets. The method also significantly improves the speed of training. This makes model combination practical, even for deep neural nets. The technique seems to reduce node interactions, leading them to learn more robust features that better generalize to new data.

(2) DropConnect

DropConnect is the generalization of dropout in which each connection, rather than each output unit, can be dropped with probability $1 - p$. Each unit thus receives input from a random subset of units in the previous layer.

DropConnect is similar to dropout as it introduces dynamic sparsity within the model, but differs in that the sparsity is on the weights, rather than the output vectors of a layer. In other words, the fully connected layer with DropConnect becomes a sparsely connected layer in which the connections are chosen at random during the training stage.

(3) Stochastic pooling

A major drawback to Dropout is that it does not have the same benefits for convolutional layers, where the neurons are not fully connected.

In stochastic pooling, the conventional deterministic pooling operations are replaced with a stochastic procedure, where the activation within each pooling region is picked randomly according to a multinomial distribution, given by the activities within the pooling region. The approach is hyperparameter free and can be combined with other regularization approaches, such as dropout and data augmentation.

An alternate view of stochastic pooling is that it is equivalent to standard max pooling but with many copies of an input image, each having small local deformations. This is similar to explicit elastic deformations of the input images, which delivers excellent MNIST performance. Using stochastic pooling

Solution (cont.)

in a multilayer model gives an exponential number of deformations since the selections in higher layers are independent of those below.

(4) Artificial data

Since the degree of model overfitting is determined by both its power and the amount of training it receives, providing a convolutional network with more training examples can reduce overfitting. Since these networks are usually trained with all available data, one approach is to either generate new data from scratch (if possible) or perturb existing data to create new ones. For example, input images could be asymmetrically cropped by a few percent to create new examples with the same label as the original.

Explicit

(1) Early stopping

One of the simplest methods to prevent overfitting of a network is to simply stop the training before overfitting has had a chance to occur. It comes with the disadvantage that the learning process is halted.

(2) Number of parameters

Another simple way to prevent overfitting is to limit the number of parameters, typically by limiting the number of hidden units in each layer or limiting network depth. For convolutional networks, the filter size also affects the number of parameters. Limiting the number of parameters restricts the predictive power of the network directly, reducing the complexity of the function that it can perform on the data, and thus limits the amount of overfitting. This is equivalent to a “zero norm”.

(3) Weight decay

A simple form of added regularizer is weight decay, which simply adds an additional error, proportional to the sum of weights (L1 norm) or squared magnitude (L2 norm) of the weight vector, to the error at each node. The level of acceptable model complexity can be reduced by increasing the proportionality constant, thus increasing the penalty for large weight vectors.

L2 regularization is the most common form of regularization. It can be implemented by penalizing the squared magnitude of all parameters directly in the objective. The L2 regularization has the intuitive interpretation of heavily penalizing peaky weight vectors and preferring diffuse weight vectors. Due to multiplicative interactions between weights and inputs this has the appealing property of encouraging the network to use all of its inputs a little rather than some of its inputs a lot.

L1 regularization is another common form. It is possible to combine L1 with L2 regularization (this is called Elastic net regularization). The L1 regularization leads the weight vectors to become sparse during optimization. In other words, neurons with L1 regularization end up using only a sparse subset of their most important inputs and become nearly invariant to the noisy inputs.

(4) Max norm constraints

Another form of regularization is to enforce an absolute upper bound on the magnitude of the weight vector for every neuron and use projected gradient descent to enforce the constraint. In practice, this corresponds to performing the parameter update as normal, and then enforcing the constraint by clamping the weight vector \vec{w} of every neuron to satisfy $\|\vec{w}\|_2 < c$. Typical values of c are order of 3-4. Some papers report improvements[48] when using this form of regularization.

HW13

December 1, 2017

1 Question 2



1.0.1 (b)



```
In [1]: from matplotlib import pyplot as plt
import numpy as np
from mpl_toolkits.mplot3d import Axes3D
```


```
In [2]: d = 2
w = np.array([
    [-1,-0.1],
    [-10,5],
    [-1,1],
    [1,1],
    [10,0]])
```

```
In [3]: for i in range(len(w)):
    fig = plt.figure()
    ax = Axes3D(fig)
    X = np.arange(0, 1, 0.05)
    Y = np.arange(0, 1, 0.05)
    X, Y = np.meshgrid(X, Y)
    Z = np.cos(X*w[i,0]+Y*w[i,1])/2/np.linalg.norm(w[i])

    ax.plot_surface(X, Y, Z, rstride=1, cstride=1, cmap='rainbow')
    plt.title(r'w=(%.1f,%.1f)'%(w[i,0],w[i,1]))
    #plt.title(r'$$w=(%f,%f)$$'%(w[i,0],w[i,1]))
    plt.savefig('2b'+str(i))
    plt.show()
```







2 Question 3

2.1 (a)

data_manager.py

```
In [2]: import os
import numpy as np
import copy
import glob
import pickle
import IPython
from matplotlib import pyplot as plt
from PIL import Image
class data_manager(object):
    def __init__(self, classes, image_size, compute_features = None, compute_label = None):

        #Batch Size for training
        self.batch_size = 40
        #Batch size for test, more samples to increase accuracy
        self.val_batch_size = 400

        self.classes = classes
```

```

self.num_class = len(self.classes)
self.image_size = image_size

self.class_to_ind = dict(zip(self.classes, range(len(self.classes))))

self.cursor = 0
self.t_cursor = 0
self.epoch = 1

self.recent_batch = []

if compute_features == None:
    self.compute_feature = self.compute_features_baseline
else:
    self.compute_feature = compute_features

if compute_label == None:
    self.compute_label = self.compute_label_baseline
else:
    self.compute_label = compute_label

self.load_train_set()
self.load_validation_set()

def get_train_batch(self):
    '''
    Compute a training batch for the neural network
    The batch size should be size 40
    '''
    batch = np.random.choice(self.train_data, self.batch_size)
    features = np.array([i['features'] for i in batch])
    labels = np.array([i['label'] for i in batch])
    return features, labels

def get_empty_state(self):
    images = np.zeros((self.batch_size, self.image_size, self.image_size, 3))
    return images

def get_empty_label(self):
    labels = np.zeros((self.batch_size, self.num_class))
    return labels

```

```

def get_empty_state_val(self):
    images = np.zeros((self.val_batch_size, self.image_size, self.image_size, 3))
    return images

def get_empty_label_val(self):
    labels = np.zeros((self.val_batch_size, self.num_class))
    return labels

def get_validation_batch(self):
    '''
    Compute a training batch for the neural network

    The batch size should be size 400

    '''
    #FILL IN
    batch = np.random.choice(self.val_data, self.val_batch_size)
    features = np.array([i['features'] for i in batch])
    labels = np.array([i['label'] for i in batch])
    return features, labels

def compute_features_baseline(self, image):
    '''
    computes the featurized on the images. In this case this corresponds
    to rescaling and standardizing.
    '''

    image = image.resize((self.image_size, self.image_size))
    image = (np.array(image) / 255.0) * 2.0 - 1.0

    return image

def compute_label_baseline(self, label):
    '''
    Compute one-hot labels given the class size
    '''

    one_hot = np.zeros(self.num_class)

    idx = self.classes.index(label)

    one_hot[idx] = 1.0

```

```

        return one_hot

def load_set(self, set_name):
    '''
    Given a string which is either 'val' or 'train', the function should load all the
    data into an
    '''

    data = []
    data_paths = glob.glob(set_name+'/*.png')

    count = 0

    for datum_path in data_paths:

        label_idx = datum_path.find('_')

        label = datum_path[len(set_name)+1:label_idx]

        if self.classes.count(label) > 0:

            img = Image.open(datum_path)

            label_vec = self.compute_label(label)

            features = self.compute_feature(img)

            data.append({'c_img': np.array(img), 'label': label_vec, 'features': fea

    np.random.shuffle(data)
    return data

def load_train_set(self):
    '''
    Loads the train set
    '''

    self.train_data = self.load_set('train')

def load_validation_set(self):

```

```

        '''
        Loads the validation set
        '''

        self.val_data = self.load_set('val')

In [3]: import numpy as np
import tensorflow as tf
#import yolo.config_card as cfg

import IPython

slim = tf.contrib.slim

class CNN(object):

    def __init__(self, classes, image_size):
        '''
        Initializes the size of the network
        '''

        self.classes = classes
        self.num_class = len(self.classes)
        self.image_size = image_size

        self.output_size = self.num_class
        self.batch_size = 40

        self.images = tf.placeholder(tf.float32, [None, self.image_size, self.image_size,

        self.logits = self.build_network(self.images, num_outputs=self.output_size)

        self.labels = tf.placeholder(tf.float32, [None, self.num_class])

        self.loss_layer(self.logits, self.labels)
        self.total_loss = tf.losses.get_total_loss()
        tf.summary.scalar('total_loss', self.total_loss)

    def build_network(self,
                      images,
                      num_outputs,
                      scope='yolo'):

        with tf.variable_scope(scope):
            with slim.arg_scope([slim.conv2d, slim.fully_connected],
                                weights_initializer=tf.truncated_normal_initializer(0.0,

```

```

        weights_regularizer=slim.l2_regularizer(0.0005)):

    '''
    Fill in network architecture here
    Network should start out with the images function
    Then it should return net
    '''

    self.conv1 = slim.conv2d(images, 5, [15, 15], activation_fn = None, scope='conv1')
    relu1 = tf.nn.relu(self.conv1)
    pool1 = slim.max_pool2d(relu1, [3,3], scope='pool1')
    fc2 = slim.fully_connected(slim.flatten(pool1), 512, activation_fn = None, scope='fc2')
    relu2 = tf.nn.relu(fc2)
    net = slim.fully_connected(relu2, 25, activation_fn = None, scope='fc3')

    return net

def get_acc(self, y_, y_out):

    '''
    Fill in a way to compute accuracy given two tensorflow arrays
    y_ (the true label) and y_out (the predict label)
    '''

    cp = tf.equal(tf.argmax(y_out, 1), tf.argmax(y_, 1))

    ac = tf.reduce_mean(tf.cast(cp, tf.float32))

    return ac

def loss_layer(self, predicts, classes, scope='loss_layer'):

    '''
    The loss layer of the network, which is written for you.
    You need to fill in get_accuracy to report the performance
    '''

    with tf.variable_scope(scope):

        self.class_loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(labels=classes, logits=predicts))

        self.accuracy = self.get_acc(classes, predicts)

```

confusion_mat.py

```

In [4]: from sklearn.metrics import confusion_matrix
        class Confusion_Matrix(object):

```

```

def __init__(self, val_data, train_data, class_labels, sess):

    self.val_data = val_data
    self.train_data = train_data
    self.CLASS_LABELS = class_labels
    self.sess = sess


def test_net(self, net):

    true_labels = []
    predicted_labels = []
    error = []
    for datum in self.val_data:

        batch_eval = np.zeros([1, datum['features'].shape[0], datum['features'].shape[1], datum['features'].shape[2], datum['features'].shape[3]])
        batch_eval[0, :, :, :] = datum['features']
        batch_label = np.zeros([1, len(self.CLASS_LABELS)])

        batch_label[0, :] = datum['label']

        prediction = self.sess.run(net.logits,
                                    feed_dict={net.images: batch_eval})

        softmax_error = self.sess.run(net.class_loss,
                                       feed_dict={net.images: batch_eval, net.labels: batch_label})

        error.append(softmax_error)

        class_pred = np.argmax(prediction)
        class_truth = np.argmax(datum['label'])

        true_labels.append(class_truth)
        predicted_labels.append(class_pred)

    self.getConfusionMatrixPlot(true_labels, predicted_labels, self.CLASS_LABELS)


def vizualize_features(self, net):

    for datum in self.val_data:

        batch_eval = np.zeros([1, datum['features'].shape[0], datum['features'].shape[1], datum['features'].shape[2], datum['features'].shape[3]])
        batch_eval[0, :, :, :] = datum['features']

        batch_label = np.zeros([1, len(self.CLASS_LABELS)])

```



```

        batch_label[0,:] = datum['label']

        response_map = self.sess.run(net.response_map,
                                      feed_dict={net.images: batch_eval, net.labels: batch_label})

        for i in range(5):
            img = self.revert_image(response_map[0,:,:,:i])
            cv2.imshow('debug',img)
            cv2.waitKey(300)

def revert_image(self,img):
    img = (img+1.0)/2.0*255.0

    img = np.array(img,dtype=int)

    blank_img = np.zeros([img.shape[0],img.shape[1],3])

    blank_img[:, :, 0] = img
    blank_img[:, :, 1] = img
    blank_img[:, :, 2] = img

    img = blank_img.astype("uint8")

    return img

def getConfusionMatrix(self,true_labels, predicted_labels):
    """
    Input
    true_labels: actual labels
    predicted_labels: model's predicted labels

    Output
    cm: confusion matrix (true labels vs. predicted labels)
    """

    # Generate confusion matrix using sklearn.metrics
    cm = confusion_matrix(true_labels, predicted_labels)
    return cm

def plotConfusionMatrix(self,cm, alphabet):
    """
    Input
    cm: confusion matrix (true labels vs. predicted labels)

```

alphabet: names of class labels

Output

Plot confusion matrix (true labels vs. predicted labels)
"""

```
fig = plt.figure()
plt.clf()                                # Clear plot
ax = fig.add_subplot(111)                # Add 1x1 grid, first subplot
ax.set_aspect(1)
res = ax.imshow(cm, cmap=plt.cm.binary,
                 interpolation='nearest', vmin=0, vmax=80)

plt.colorbar(res)                        # Add color bar

width = len(cm)                          # Width of confusion matrix
height = len(cm[0])                     # Height of confusion matrix

# Annotate confusion entry with numeric value

for x in range(width):
    for y in range(height):

        ax.annotate(str(cm[x][y]), xy=(y, x), horizontalalignment='center',
                    verticalalignment='center', color=self.get_font_color(cm[x][y]))

# Plot confusion matrix (true labels vs. predicted labels)
plt.xticks(range(width), alphabet[:width], rotation=90)
plt.yticks(range(height), alphabet[:height])
plt.show()
return plt

def getConfusionMatrixPlot(self, true_labels, predicted_labels, alphabet):
    """
    Input
    true_labels: actual labels
    predicted_labels: model's predicted labels
    alphabet: names of class labels

    Output
    Plot confusion matrix (true labels vs. predicted labels)
    """

    # Generate confusion matrix using sklearn.metrics
    cm = confusion_matrix(true_labels, predicted_labels)
```

```

# Plot confusion matrix (true labels vs. predicted labels)
return self.plotConfusionMatrix(cm, alphabet)

```

```

def getFontColor(self,value):
    """
    Input
    value: confusion entry value

    Output
    font color for confusion entry
    """
    if value < -1:
        return "black"
    else:
        return "white"

```

test_cnn_part_a.py

```
In [26]: tf.reset_default_graph()
```

```

np.random.seed(0)

CLASS_LABELS = ['apple','banana','nectarine','plum','peach','watermelon','pear','mango',
                'radish','carrot','potato','tomato','bellpepper','broccoli','cabbage','cauliflower']

image_size = 90
classes = CLASS_LABELS
dm = data_manager(classes, image_size)


cnn = CNN(classes,image_size)

sess = tf.Session()
sess.run(tf.global_variables_initializer())

val_data = dm.val_data
train_data = dm.train_data

cm = Confusion_Matrix(val_data,train_data,CLASS_LABELS,sess)
cm.test_net(cnn)

```



2.2 (c)

trainer.py

```
In [5]: import datetime
import os
import sys
import argparse

class Solver(object):

    def __init__(self, net, data):

        self.net = net
        self.data = data

        self.max_iter = 5000
        self.summary_iter = 200

        self.learning_rate = 0.1
```

```

self.saver = tf.train.Saver()

self.summary_op = tf.summary.merge_all()

self.global_step = tf.get_variable(
    'global_step', [], initializer=tf.constant_initializer(0), trainable=False)

'''
Tensorflow is told to use a gradient descent optimizer
In the function optimize you will iteratively apply this on batches of data
'''

self.train_step = tf.train.MomentumOptimizer(.003, .9)
self.train = self.train_step.minimize(self.net.class_loss)


self.sess = tf.Session()
self.sess.run(tf.global_variables_initializer())


def optimize(self):

    self.train_losses = []
    self.test_losses = []
    self.train_accuracy = []
    self.test_accuracy = []
    '''
    Performs the training of the network.
    Implement SGD using the data manager to compute the batches
    Make sure to record the training and test loss through out the process
    '''
    for i in range(self.max_iter):
        x,y = self.data.get_train_batch()
        _,loss,accuracy = self.sess.run([self.train,
                                         self.net.class_loss,
                                         self.net.accuracy],
                                         feed_dict={self.net.images: x, self.net.labels: y})
        self.train_losses.append(loss)
        self.train_accuracy.append(accuracy)

    for i in range(self.summary_iter):
        x,y = self.data.get_validation_batch()
        loss,accuracy = self.sess.run([self.net.class_loss,
                                       self.net.accuracy],
                                       feed_dict={self.net.images: x, self.net.labels: y})
        self.test_losses.append(loss)
        self.test_accuracy.append(accuracy)

```

viz_features.py

```
In [6]: from sklearn.metrics import confusion_matrix

class Viz_Feat(object):

    def __init__(self, val_data, train_data, class_labels, sess):

        self.val_data = val_data
        self.train_data = train_data
        self.CLASS_LABELS = class_labels
        self.sess = sess

    def vizualize_features(self, net):

        images = [0, 10, 100]
        '''
        Compute the response map for the index images
        '''
        for j in images:
            batch_eval = np.zeros([1, self.train_data[0]['features'].shape[0],
                                   self.train_data[0]['features'].shape[1],
                                   self.train_data[0]['features'].shape[2]])
            batch_eval[0, :, :, :] = self.train_data[j]['features']

            response_map = self.sess.run(net.conv1,
                                          feed_dict={net.images: batch_eval})

            for i in range(5):
                plt.subplot(1, 5, i+1)
                img = self.revert_image(response_map[0, :, :, i])
                plt.imshow(img)
            plt.show()

    def revert_image(self, img):
        '''
        Used to revert images back to a form that can be easily visualized
        '''

        img = (img+1.0)/2.0*255.0

        img = np.array(img, dtype=int)

        blank_img = np.zeros([img.shape[0], img.shape[1], 3])

        blank_img[:, :, 0] = img
        blank_img[:, :, 1] = img
```

```

blank_img[:, :, 2] = img

img = blank_img.astype("uint8")

return img

```

train_cnn.py

```

In [29]: tf.reset_default_graph()
        CLASS_LABELS = ['apple', 'banana', 'nectarine', 'plum', 'peach', 'watermelon', 'pear', 'mango',
                        'radish', 'carrot', 'potato', 'tomato', 'bellpepper', 'broccoli', 'cabbage', 'cauliflower']

        LITTLE_CLASS_LABELS = ['apple', 'banana', 'eggplant']

        image_size = 90

        np.random.seed(0)

        classes = CLASS_LABELS
        dm = data_manager(classes, image_size)
        with tf.variable_scope("c"):
            cnn = CNN(classes, image_size)

            solver = Solver(cnn, dm)



            solver.optimize()

        plt.plot(solver.test_losses, label = 'Validation')
        plt.legend()
        plt.xlabel('Iterations')
        plt.ylabel('Loss')
        plt.show()
        plt.plot(solver.train_losses, label = 'Training')
        plt.legend()
        plt.xlabel('Iterations')
        plt.ylabel('Loss')
        plt.show()

        val_data = dm.val_data
        train_data = dm.train_data

        sess = solver.sess




```




```
In [30]: cm = Viz_Feat(val_data,train_data,CLASS_LABELS,sess)
```




```
cm.vizualize_features(cnn)
```



```
In [31]: plt.plot(solver.train_accuracy,label = 'Training')  
plt.legend()  
plt.xlabel('Iterations')  
plt.ylabel('Accuracy')  
plt.show()
```



```
In [32]: plt.plot(solver.test_accuracy,label = 'Validation')
plt.legend()
plt.xlabel('Iterations (in 200s)')
plt.ylabel('Accuracy')
plt.show()
```



nn_classifier.py

```
In [7]: from numpy.random import uniform
import time
import glob
from sklearn.neighbors import KNeighborsClassifier

class NN():
    def __init__(self, train_data, val_data, n_neighbors=5):

        self.train_data = train_data
        self.val_data = val_data

        self.sample_size = 400
        self.model = KNeighborsClassifier(n_neighbors=n_neighbors)

    def train_model(self):

        '''
        Train Nearest Neighbors model
        '''
        x = np.array([np.reshape(self.train_data[i]['features'], (90*90*3)) for i in range(len(self.train_data))])
        y = np.array([self.train_data[i]['label'] for i in range(len(self.train_data))])
        self.model.fit(x, y)
```

```

def get_validation_error(self):
    '''
    Compute validation error. Please only compute the error on the sample_size number
    over randomly selected data points. To save computation.
    '''
    index = np.random.choice(len(self.val_data),self.sample_size, replace=False)
    predy = self.model.predict([np.reshape(self.val_data[i]['features'],(90*90*3)) for i in index])
    return np.mean(np.argmax(predy,1) != np.argmax([self.val_data[i]['label'] for i in index]))

def get_train_error(self):
    '''
    Compute train error. Please only compute the error on the sample_size number
    over randomly selected data points. To save computation.
    '''
    index = np.random.choice(len(self.train_data),self.sample_size, replace=False)
    predy = self.model.predict([np.reshape(self.train_data[i]['features'],(90*90*3)) for i in index])
    return np.mean(np.argmax(predy,1) != np.argmax([self.train_data[i]['label'] for i in index]))

```

train_nn.py

```

In [ ]: CLASS_LABELS = ['apple','banana','nectarine','plum','peach','watermelon','pear','mango',
    'radish','carrot','potato','tomato','bellpepper','broccoli','cabbage','cauliflower',

image_size = 90
classes = CLASS_LABELS

dm = data_manager(classes, image_size)

val_data = dm.val_data
train_data = dm.train_data

K = [1, 20, 100]
test_losses = []
train_losses = []

for k in K:
    nn = NN(val_data,train_data,n_neighbors=k)


    nn.train_model()

    test_losses.append(nn.get_validation_error())
    train_losses.append(nn.get_train_error())


In [9]: plt.plot(K, 1-np.array(test_losses),label = 'Validation')

```

```
plt.plot(K, 1-np.array(train_losses), label = 'Training')
plt.legend()
plt.xlabel('Number of Neighbors')
plt.ylabel('Accuracy')
plt.show()
```



```
In [11]: plt.plot(K, test_losses, label = 'Validation')
plt.plot(K, train_losses, label = 'Training')
plt.legend()
plt.xlabel('Number of Neighbors')
plt.ylabel('Error')
plt.show()
```



In []: