

**STAT 309: MATHEMATICAL COMPUTATIONS I**  
**FALL 2019**  
**LECTURE 0**

1. LAUNDRY LIST

- web site: <http://www.stat.uchicago.edu/~lekheng/courses/309/>
- notes: <http://www.stat.uchicago.edu/~lekheng/courses/309/notes/>
- books: <http://www.stat.uchicago.edu/~lekheng/courses/309/books/>
- last year's notes: <http://www.stat.uchicago.edu/~lekheng/courses/309f18/notes/>
- no required textbook
- useful references:
  - Trefethen and Bau
  - Watkins
  - Golub and Van Loan
  - Demmel
- facts about matrices:
  - Bernstein
- homework:
  - homework due beginning of class
  - collaboration allowed but must be declared
  - six assignments, lowest score will be dropped, accounting for 50% of grade
  - no late homework will be accepted
- exams:
  - two 80-minute quizzes or one 3-hour exam (to be decided later)
  - in-class, closed-book, no cheat sheet, 1 hr 20 min
- grade: 50% homework, 50% quizzes/exam

2. NUMERICAL ANALYSIS

- numerical analysis: study of algorithms for continuous mathematics
- examples:
  - linear partial differential equation: given  $c_\alpha$ 's, find  $f$

$$\sum_{|\alpha| \leq n} c_\alpha(\mathbf{t}) \frac{\partial^\alpha}{\partial \mathbf{t}^\alpha} f(\mathbf{t}) = 0 \quad (2.1)$$

- Fredholm integral equation of the first kind: given  $K$  and  $g$ , find  $f$

$$\int_{\Omega} K(\mathbf{s}, \mathbf{t}) f(\mathbf{t}) d\mathbf{t} = g(\mathbf{s}) \quad (2.2)$$

- linear eigenvalue problem: given  $c_\alpha$ 's, find  $f$  and  $\lambda$

$$\sum_{|\alpha| \leq n} c_\alpha(\mathbf{t}) \frac{\partial^\alpha}{\partial \mathbf{t}^\alpha} f(\mathbf{t}) = \lambda f(\mathbf{t}) \quad (2.3)$$

- Fredholm integral equation of the first kind: given  $K$  and  $g$ , find  $f$  and  $\lambda$

$$g(\mathbf{s}) + \lambda \int_{\Omega} K(\mathbf{s}, \mathbf{t}) f(\mathbf{t}) d\mathbf{t} = f(\mathbf{s}) \quad (2.4)$$

- nonlinear optimization: given  $f_0, \dots, f_m$ , find  $\mathbf{t}_{\min}$

$$\min f_0(\mathbf{t}) \quad \text{subject to} \quad f_1(\mathbf{t}) \leq 0, \dots, f_m(\mathbf{t}) \leq 0 \quad (2.5)$$

- many scientific and engineering can be formulated in one of these forms — the PDE or integral equations would be a mathematical formulation of physical principles like Newton's second law or Maxwell equations or Schrödinger equation
- we can rarely solve these analytically, i.e., give a useful closed-form formula for the solution
- have to rely on computers, which can only deal with discrete problems
- discretization of (2.1) or (2.2), or Newton method applied to (2.5) yields

$$A\mathbf{x} = \mathbf{b} \quad (2.6)$$

- discretization of (2.3) or (2.4) yields

$$A\mathbf{x} = \lambda\mathbf{x} \quad (2.7)$$

- when we discretize, we have

$$\mathbf{x} = \begin{bmatrix} f(\mathbf{t}_1) \\ f(\mathbf{t}_2) \\ \vdots \\ f(\mathbf{t}_n) \end{bmatrix}$$

- solving for  $\mathbf{x}$  gives us a sample of point values of  $f$ , which is often enough for many purposes
- the larger  $n$  is, the more information we get about  $f$
- the matrix  $A$  comes from discretization of the linear operator — differential operators in the case of (2.1) or (2.3) and integral operators in the case of (2.2) or (2.4)
- example: discretizing a 1-dimensional differential operator

$$\frac{d^2}{dt^2} \xrightarrow{\text{discretize}} \begin{bmatrix} 2 & -1 & & \\ -1 & 2 & \ddots & \\ & \ddots & \ddots & -1 \\ & & -1 & 2 \end{bmatrix} \in \mathbb{R}^{n \times n}$$

- example: discretizing a 2-dimensional differential operator

$$\frac{\partial^2}{\partial t_1^2} + \frac{\partial^2}{\partial t_2^2} \xrightarrow{\text{discretize}} \begin{bmatrix} D & -I & & \\ -I & D & \ddots & \\ & \ddots & \ddots & -I \\ & & -I & D \end{bmatrix} \in \mathbb{R}^{mn \times mn}, \text{ where } D = \begin{bmatrix} 4 & -1 & & \\ -1 & 4 & \ddots & \\ & \ddots & \ddots & -1 \\ & & -1 & 4 \end{bmatrix} \in \mathbb{R}^{m \times m}$$

- bottom line: many problems in science and engineering require that we solve (2.6) or (2.7)

### 3. OPTIMIZATION

- suppose you want to solve an optimization problem

$$\begin{array}{ll} \text{minimize} & f(\mathbf{x}) \\ \text{subject to} & h_i(\mathbf{x}) \leq 0, \quad i = 1, \dots, p, \\ & A\mathbf{x} = \mathbf{b} \end{array}$$

- one of the most widely used algorithm is interior point method (essentially Newton's method adapted to a constrained optimization problem) which requires us to solve a linear system of the form

$$\begin{bmatrix} t\nabla^2 f(\mathbf{x}_k) + \nabla^2 \varphi(\mathbf{x}_k) & A^\top \\ A & 0 \end{bmatrix} \begin{bmatrix} \Delta \mathbf{x}_k \\ \nu_k \end{bmatrix} = - \begin{bmatrix} t\nabla f(\mathbf{x}_k) + \nabla \varphi(\mathbf{x}_k) \\ 0 \end{bmatrix}$$

where  $\varphi$  is the so-called log barrier function that 'traps' the iterates  $\mathbf{x}_k$  within the region defined by the constraints

- at each iterate  $\mathbf{x}_k$ , we will have to solve such a linear system for  $\Delta \mathbf{x}_k$  to obtain the next iterate  $\mathbf{x}_{k+1} = \mathbf{x}_k + \Delta \mathbf{x}_k$
- so the computational cost of interior point methods is largely dominated by the cost of solving linear systems

#### 4. MACHINE LEARNING

- many modern problems are information theoretic in nature
  - no differential or integral equations describing your solution  $f$
  - but a large *test set* of given data  $\{(x_i, f(x_i)) : i = 1, \dots, n\}$  that allows you to guess your  $f$
- example: classification problems
  - spam identification

$$f : \text{emails} \rightarrow \{\text{spam}, \text{nospam}\}$$

- image recognition

$$f : \text{facial images} \rightarrow \{\text{male}, \text{female}\}$$

or more generally

$$f : \text{handwritten digits} \rightarrow \{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$$

- there is no 'Newton's law' type of rule to describe  $f$
- example: supervised learning for binary classification

$$f : X \rightarrow \{-1, +1\}$$

- given training set  $\Omega = \{x_1, \dots, x_n\} \subseteq X$ , i.e., we already know the value  $f(x_i) = y_i$  for any  $x_i \in \Omega$
- want to find  $f$ , i.e., given some  $x \notin \Omega$ , we want to predict the value  $f(x)$
- let us use spam identification as an example, then for any e-mail  $x \in X$ ,

$$f(x) = \begin{cases} -1 & \text{if } x \text{ is spam} \\ +1 & \text{if } x \text{ is not spam} \end{cases}$$

- we can encode an e-mail as a vector in  $\mathbb{R}^N$ , for example, by counting word frequencies
- so if you like you may assume that  $X \subseteq \mathbb{R}^N$  where  $N$  is very large
- one way to do this:
  - assume that

$$f(x) = \sum_{i=1}^n c_i K(x, x_i) \tag{4.1}$$

where  $K : X \times X \rightarrow \mathbb{R}$  is some suitable *Mercer kernel*

- if  $X \subseteq \mathbb{R}^N$  a common example is the Gaussian kernel

$$K(\mathbf{x}, \mathbf{y}) = e^{-\|\mathbf{x} - \mathbf{y}\|^2 / 2\sigma^2}$$

- since we already know the value of  $f(x)$  for  $x \in \{x_1, \dots, x_n\}$ , we could in principle determine  $c_1, \dots, c_n$  by plugging  $x_1, \dots, x_n$  into (4.1) to get

$$\begin{aligned} f(x_1) &= c_1 K(x_1, x_1) + \dots + c_n K(x_1, x_n) \\ f(x_2) &= c_1 K(x_2, x_1) + \dots + c_n K(x_2, x_n) \\ &\vdots \\ f(x_n) &= c_1 K(x_n, x_1) + \dots + c_n K(x_n, x_n) \end{aligned}$$

or equivalently

$$\begin{bmatrix} K(x_1, x_1) & K(x_1, x_2) & \dots & K(x_1, x_n) \\ K(x_2, x_1) & \ddots & & \vdots \\ \vdots & & \ddots & \vdots \\ K(x_n, x_1) & \dots & \dots & K(x_n, x_n) \end{bmatrix} \begin{bmatrix} c_1 \\ c_2 \\ \vdots \\ c_n \end{bmatrix} = \begin{bmatrix} f(x_1) \\ f(x_2) \\ \vdots \\ f(x_n) \end{bmatrix} \quad (4.2)$$

or

$$K\mathbf{c} = \mathbf{y}$$

- note that we know all values  $K(x_i, x_j)$  of the matrix and also the right-hand side  $f(x_i)$  as long as we have the training set  $\{(x_i, f(x_i)) : i = 1, \dots, n\}$
- so we end up with a linear system like (2.6) again
- in principle this is very nice but in practice it rarely works since (4.2) is unlikely to have a solution
- so what we often need to do is to solve linear systems (2.6) approximately, i.e.,  $A\mathbf{x} \approx \mathbf{b}$  where ‘ $\approx$ ’ is interpreted in some appropriate ways — we will look at some of these variants of (2.6) later
- the most common interpretation of  $A\mathbf{x} \approx \mathbf{b}$  is the *least squares problem*

$$\min_{\mathbf{x} \in \mathbb{R}^n} \|A\mathbf{x} - \mathbf{b}\|_2^2 = \min_{x_1, \dots, x_n \in \mathbb{R}} \sum_{i=1}^m \sum_{j=1}^n (a_{ij}x_j - b_i)^2 \quad (4.3)$$

- in the context of supervised learning, this is called *empirical risk minimization*, i.e., find  $c_1, \dots, c_n$  so that

$$\sum_{i=1}^n (y_i - f(x_i))^2 \quad (4.4)$$

is minimized

- but (4.4) is often *ill-posed* (no unique solution) and so a common strategy is to do Tikhonov regularization and minimize instead

$$\sum_{i=1}^n (y_i - f(x_i))^2 + \lambda \|f\|_K^2$$

where  $\|\cdot\|_K$  is a special norm induced by the kernel  $K(x, y)$  (if you must know, it is called the reproducing kernel Hilbert space or RKHS norm)

- as we will see later in this course, this leads to a problem of the form

$$(K + \lambda I)\mathbf{c} = \mathbf{y}$$

which is again a linear system except that we will need to find  $\lambda$  separately (we will see how to do this)

- now once we have  $c_1, \dots, c_n$ , given any  $x$ , we can find the value  $f(x)$

- of course  $f(x)$  would in general not be  $\pm 1$  but we can design a rule of the form

$$f(x) \begin{cases} < 0 & \Rightarrow x \text{ is spam} \\ > 0 & \Rightarrow x \text{ is not spam} \end{cases}$$

- so we have built a spam filter

## 5. SOLVING LINEAR SYSTEMS

- most of the course will focus on solving *linear systems* (2.6) and its variants like least squares, regularized least squares, total least squares, etc
- the fundamental problem is

$$A\mathbf{x} = \mathbf{b}$$

where we are given  $A \in \mathbb{C}^{m \times n}$ ,  $\mathbf{b} \in \mathbb{C}^m$  and we seek a solution  $\mathbf{x} \in \mathbb{C}^n$

- often we will work over  $\mathbb{R}$  instead of  $\mathbb{C}$  but these would be only fields of interest
- some of the stuff we say in this course will be false over arbitrary fields (e.g.  $\mathbb{F}_2 = \{0, 1\}$  with mod 2 arithmetic)
- three important numbers associated to a matrix  $A$  or a linear system  $A\mathbf{x} = \mathbf{b}$ :
  - $m$  = number of rows = number of equations
  - $n$  = number of columns = number of variables
  - $r = \text{rank}(A) = \dim(\text{im}(A)) = \dim(\text{colsp}(A)) = \dim(\text{rowsp}(A))$
- $m, n, r$  tell us about existence and uniqueness of solution to  $A\mathbf{x} = \mathbf{b}$
- terminologies
  - $m = n$ :  $A$  is square matrix,  $A\mathbf{x} = \mathbf{b}$  is a square system, i.e. number of variables equals number of equations

$$A = \begin{bmatrix} \times & \times & \times & \times \\ \times & \times & \times & \times \\ \times & \times & \times & \times \\ \times & \times & \times & \times \end{bmatrix}$$

- $m > n$ :  $A$  is tall-and-thin matrix,  $A\mathbf{x} = \mathbf{b}$  is an overdetermined system, i.e., more equations than variables

$$A = \begin{bmatrix} \times & \times \\ \times & \times \\ \times & \times \\ \times & \times \end{bmatrix}$$

- $m < n$ :  $A$  is short-and-fat matrix,  $A\mathbf{x} = \mathbf{b}$  is an underdetermined system, i.e., more variables than equations

$$A = \begin{bmatrix} \times & \times & \times & \times \\ \times & \times & \times & \times \end{bmatrix}$$

- if  $r = \min\{m, n\}$ , we say  $A$  is of *full rank*, otherwise we say  $A$  is *rank deficient*
- if  $r = \min\{m, n\} = m$ , we say  $A$  is of *full row rank*
- if  $r = \min\{m, n\} = n$ , we say  $A$  is of *full column rank*
- question: what is the big deal about solving linear systems  $A\mathbf{x} = \mathbf{b}$ ? don't we know all about this already?
- answer: we only know how to solve idealized versions of the problem, but not in realistic situations
  - what if there are rounding errors in the coefficient matrix  $A$  or the right hand side  $\mathbf{b}$  or both
  - what if we want to solve it quicker than  $O(n^3)$
  - what if  $m$  and  $n$  are large

- what if we want to do things in parallel on multicore processors
- what if we need to deal with a variant with constraints on the solution  $\mathbf{x}$ , or where  $A\mathbf{x} = \mathbf{b}$  has no solution or no unique solution (as we saw in the machine learning example)
- linear systems are arguably the most widely solved problem in science and engineering
  - 70% of supercomputing time is spent on this
  - that’s why solution of linear system is used to benchmark supercomputers (cf. <http://www.top500.org>)

## 6. TOP 10 ALGORITHMS OF THE 20TH CENTURY

- a broader motivation for this course and its sequel next quarter is that matrix computations are behind some of the most important algorithms
  - the three bold faced ones are algorithms in matrix computations
  - the four italics ones are algorithms are variants or extensions of algorithms in matrix computations
  - see <http://www.stat.uchicago.edu/~lekheng/courses/309/top10/>
- (1) Metropolis Algorithm for Monte Carlo
  - (2) *Simplex Method for Linear Programming*
  - (3) **Krylov Subspace Iteration Methods**
  - (4) **Decompositional Approach to Matrix Computations**
  - (5) Fortran Optimizing Compiler
  - (6) **QR Algorithm for Computing Eigenvalues**
  - (7) Quicksort Algorithm for Sorting
  - (8) *Fast Fourier Transform*
  - (9) *Integer Relation Detection*
  - (10) *Fast Multipole Method*

## 7. VARIANTS OF $A\mathbf{x} = \mathbf{b}$

- notations
  - $\mathbf{x} = [x_1, \dots, x_n]^T \in \mathbb{R}^n$

$$\|\mathbf{x}\|_2 := \sqrt{x_1^2 + \dots + x_n^2}$$

$$\|\mathbf{x}\|_1 := |x_1| + \dots + |x_n|$$

$$\|\mathbf{x}\|_\infty := \max\{|x_1|, \dots, |x_n|\}$$

$$\|\mathbf{x}\|_0 := \text{nnz}(\mathbf{x}) = \#\{i : x_i \neq 0\}$$

$$- A = [a_{ij}]_{i,j=1}^{m,n} \in \mathbb{R}^{m \times n}$$

$$\|A\|_F = \sqrt{\sum_{i,j=1}^{m,n} |a_{ij}|^2}$$

- we will discuss vector and matrix norms below
- note that  $\|\cdot\|_0$  is not a norm

- (1) linear regression or least squares problem: know  $A$  exactly but  $\mathbf{b}$  is corrupted by error  $\mathbf{r}$ , i.e.,  $A\mathbf{x} = \mathbf{b} + \mathbf{r}$ , and we want an  $\mathbf{x}$  that minimizes  $\mathbf{r}$ ,

$$\min\{\|\mathbf{r}\|_2^2 : A\mathbf{x} = \mathbf{b} + \mathbf{r}\} = \min_{\mathbf{x} \in \mathbb{R}^n} \|A\mathbf{x} - \mathbf{b}\|_2^2 \quad (7.1)$$

Gauss–Markov theorem says that such an  $\mathbf{x}$  is the maximum likelihood estimator if the error  $\mathbf{r}$  is from a distribution that has zero mean and finite variance

- (2) error-in-variables regression or total least squares problem:  $A$  and  $\mathbf{b}$  are both corrupted by error  $E$  and  $\mathbf{r}$ , i.e.,  $(A + E)\mathbf{x} = \mathbf{b} + \mathbf{r}$ , and we want an  $\mathbf{x}$  that minimizes both  $E$  and  $\mathbf{r}$ ,

$$\min\{\|E\|_F^2 + \|\mathbf{r}\|_2^2 : (A + E)\mathbf{x} = \mathbf{b} + \mathbf{r}\}$$

- (3) data least squares problem:  $A$  is corrupted by error  $E$ , i.e.,  $(A + E)\mathbf{x} = \mathbf{b}$ , and we want an  $\mathbf{x}$  that minimizes  $E$ ,

$$\min\{\|E\|_F^2 : (A + E)\mathbf{x} = \mathbf{b}\}$$

- (4) minimum norm least squares: want the minimum length solution to (7.1),

$$\min\{\|\mathbf{x}\|_2^2 : \mathbf{x} \in \operatorname{argmin}\|A\mathbf{x} - \mathbf{b}\|_2^2\} = \min\{\|\mathbf{x}\|_2^2 : A^\top A\mathbf{x} = A^\top \mathbf{b}\} \quad (7.2)$$

the solution  $\mathbf{x}_*$  to (7.2) is unique and can in fact be used to define the Moore–Penrose pseudoinverse of  $A$ :  $\mathbf{x}_* = A^\dagger \mathbf{b}$

- (5) robust regression: replace 2-norm by 1-norm (more generally, the Huber loss function) in (7.1),

$$\min\{\|\mathbf{r}\|_1 : A\mathbf{x} = \mathbf{b} + \mathbf{r}\} = \min_{\mathbf{x} \in \mathbb{R}^n} \|A\mathbf{x} - \mathbf{b}\|_1$$

great for reducing sensitivity to outliers

- (6) ridge regression or regularized least squares

$$\min_{\mathbf{x} \in \mathbb{R}^n} \|A\mathbf{x} - \mathbf{b}\|_2^2 + \|\Gamma\mathbf{x}\|_2^2$$

where  $\Gamma \in \mathbb{R}^{p \times n}$  is some other matrix — most commonly  $\Gamma = \lambda I$  or the finite-difference matrix

- (7) LASSO or  $l^1$ -regularized least squares

$$\min_{\mathbf{x} \in \mathbb{R}^n} \|A\mathbf{x} - \mathbf{b}\|_2^2 + \|\Gamma\mathbf{x}\|_1$$

where  $\Gamma \in \mathbb{R}^{p \times n}$  is some other matrix — most commonly  $\Gamma = \lambda I$  or the finite-difference matrix

- (8) sparse or structured linear systems: sparse means  $A$  has a lot of zeroes (sufficiently many that it pays to take advantage of the fact), structured means that  $A$  can be defined with fewer than the usual number of  $mn$  parameters. An example of a data sparse matrix is a Toeplitz matrix

$$T = \begin{bmatrix} a_0 & a_1 & a_2 & \cdots & a_{n-1} \\ a_{-1} & a_0 & a_1 & \cdots & \\ a_{-2} & a_{-1} & a_0 & \cdots & a_1 \\ & \ddots & \ddots & \ddots & \\ a_{-n+1} & & a_{-2} & a_{-1} & a_0 \end{bmatrix} \in \mathbb{R}^{n \times n}$$

i.e.,  $a_{ij}$  depends only on  $|i - j|$ , and  $T$  can be specified with just  $2n - 1$  parameters  $a_{-n+1}, \dots, a_{n-1} \in \mathbb{R}$ ; a Toeplitz  $T\mathbf{x} = \mathbf{b}$  can be solved in  $O(n \log^2 n)$  time as opposed to the usual  $O(n^3)$  time for general linear systems

- (9) linear programming:

$$\min\{\mathbf{c}^\top \mathbf{x} : A\mathbf{x} \leq \mathbf{b}\}$$

note that  $\mathbf{c}^\top \mathbf{x} = c_1 x_1 + \dots + c_n x_n$  is a linear function; this is very important in economics

- (10) quadratic programming: given  $A \in \mathbb{R}^{n \times n}$ ,  $B \in \mathbb{R}^{m \times n}$ ,  $\mathbf{c} \in \mathbb{R}^n$ ,  $\mathbf{d} \in \mathbb{R}^m$ , want

$$\min\left\{\frac{1}{2}\mathbf{x}^\top A\mathbf{x} - \mathbf{c}^\top \mathbf{x} : B\mathbf{x} = \mathbf{d}\right\}$$

this reduces to a linear system

$$\begin{bmatrix} A & B^\top \\ B & 0 \end{bmatrix} \begin{bmatrix} \mathbf{x} \\ \boldsymbol{\lambda} \end{bmatrix} = \begin{bmatrix} \mathbf{c} \\ \mathbf{d} \end{bmatrix}$$

- (11) basis pursuit: if we want the sparsest solution to an underdetermined linear system, we ought to solve

$$\min\{\|\mathbf{x}\|_0 : A\mathbf{x} = \mathbf{b}\}$$

but this is NP-hard and so we look at a convex relaxation

$$\min\{\|\mathbf{x}\|_1 : A\mathbf{x} = \mathbf{b}\}$$

which can in fact be reduced to a linear programming problem

- we will not say much about (5) or (7) but they can be solved using least squares techniques in this course, using a method called *iteratively reweighted least squares*
- depending on whether we have time, we may or may not cover (9) and (11), but they can be solved using the *simplex algorithm* based on Gaussian elimination with partial pivoting, which we will cover

## 8. CONTINUITY OF NORMS

- all norms are continuous functions — an simple but important observation
- what can we say about the norm of the difference of two vectors? we know that  $\|\mathbf{x} - \mathbf{y}\| \leq \|\mathbf{x}\| + \|\mathbf{y}\|$  but we can obtain a more useful relationship as follows:

$$\|\mathbf{x}\| = \|(\mathbf{x} - \mathbf{y}) + \mathbf{y}\| \leq \|\mathbf{x} - \mathbf{y}\| + \|\mathbf{y}\|$$

we obtain

$$\|\mathbf{x} - \mathbf{y}\| \geq \|\mathbf{x}\| - \|\mathbf{y}\|$$

- thirdly, from

$$\|\mathbf{y}\| = \|\mathbf{y} - \mathbf{x} + \mathbf{x}\| \leq \|\mathbf{x} - \mathbf{y}\| + \|\mathbf{x}\|$$

it follows that

$$\|\mathbf{x} - \mathbf{y}\| \geq \|\mathbf{y}\| - \|\mathbf{x}\|$$

and therefore

$$\left| \|\mathbf{x}\| - \|\mathbf{y}\| \right| \leq \|\mathbf{x} - \mathbf{y}\| \tag{8.1}$$

- the inequality (8.1) yields a very important property of norms, namely, they are all (uniformly) continuous functions of the entries of their arguments — in fact, they are *Lipschitz functions* if you know what those are

## 9. EQUIVALENCE OF NORMS

- there are also interesting relationships for two different norms
- first and foremost, on *finite dimensional spaces* (which include  $\mathbb{C}^n$  and  $\mathbb{C}^{m \times n}$ ) *all norms are equivalent*
  - that is, given two norms  $\|\cdot\|_\alpha$  and  $\|\cdot\|_\beta$ , there exist constants  $c_1$  and  $c_2$  with  $0 < c_1 < c_2$  such that

$$c_1 \|\mathbf{x}\|_\alpha \leq \|\mathbf{x}\|_\beta \leq c_2 \|\mathbf{x}\|_\alpha \tag{9.1}$$

for all  $\mathbf{x} \in V$

- example: from the definition of the  $\infty$ -norm, we have

$$\|\mathbf{x}\|_\infty \leq \|\mathbf{x}\|_2 \leq \sqrt{n} \|\mathbf{x}\|_\infty$$

- example: also not hard to show that

$$\frac{1}{n} \|\mathbf{x}\|_1 \leq \|\mathbf{x}\|_\infty \leq \|\mathbf{x}\|_1$$



- in fact, no matter what crazy choices of norms that we make, say

$$\|x\|_\alpha = \left( \sum_{i=1}^n i |x_i|^n \right)^{1/n}, \quad \|x\|_\beta = \mathbf{x}^\top \begin{bmatrix} 3 & -1 & & \\ -1 & 3 & \ddots & \\ & \ddots & \ddots & -1 \\ & & -1 & 3 \end{bmatrix} \mathbf{x},$$

we know that there are  $c_1$  and  $c_2$  so that (9.1) holds

- by definition, a sequence of vectors  $\mathbf{x}_0, \mathbf{x}_1, \dots$  converges to a vector  $\mathbf{x}$  if and only if

$$\lim_{k \rightarrow \infty} \|\mathbf{x}_k - \mathbf{x}\| = 0$$

for any norm (you may also write down a formal version in terms of  $\varepsilon$  and  $N$ )

- the equivalence of norms on finite dimensional vector spaces tells us that

$$\lim_{k \rightarrow \infty} \|\mathbf{x}_k - \mathbf{x}\|_\alpha = 0 \quad \text{if and only if} \quad \lim_{k \rightarrow \infty} \|\mathbf{x}_k - \mathbf{x}\|_\beta = 0$$

for any choice of norms  $\|\cdot\|_\alpha$  and  $\|\cdot\|_\beta$  (why?)

- if we can establish convergence of an algorithm in a specific norm convergence in every other norm follows automatically
- for this reason, norms are very useful to measure the error in an approximation
- secondly we have a relationship that applies to products of norms, the *Hölder inequality*

$$|\mathbf{x}^* \mathbf{y}| \leq \|\mathbf{x}\|_p \|\mathbf{y}\|_q, \quad \frac{1}{p} + \frac{1}{q} = 1$$

- a well-known corollary arises when  $p = q = 2$ , the *Cauchy-Schwarz inequality*

$$|\mathbf{x}^* \mathbf{y}| \leq \|\mathbf{x}\|_2 \|\mathbf{y}\|_2$$

- you will see a generalization of Cauchy-Schwarz inequality called the *Bessel inequality* in Homework 0
- by setting  $\mathbf{x} = [1, 1, \dots, 1]^\top$ , the Hölder inequality yields the relationships

$$\left| \sum_{i=1}^n y_i \right| \leq \sum_{i=1}^n |y_i|$$

and

$$\left| \sum_{i=1}^n y_i \right| \leq n \max_{i=1, \dots, n} |y_i|$$

and

$$\left| \sum_{i=1}^n y_i \right| \leq \sqrt{n} \left( \sum_{i=1}^n |y_i|^2 \right)^{1/2}$$

## 10. MATRIX NORMS

- note that the space of complex  $m \times n$  matrices  $\mathbb{C}^{m \times n}$  is a vector space over  $\mathbb{C}$  (ditto for real matrices over  $\mathbb{R}$ ) of dimension  $mn$
- we write  $O$  for the  $m \times n$  zero matrix, i.e., all entries are 0
- a norm on either  $\mathbb{C}^{m \times n}$  or  $\mathbb{R}^{m \times n}$  is called a *matrix norm*
- recall that these means  $\|\cdot\| : \mathbb{C}^{m \times n} \rightarrow \mathbb{R}$  satisfies
  - (1)  $\|A\| \geq 0$  for all  $A \in \mathbb{C}^{m \times n}$
  - (2)  $\|A\| = 0$  if and only if  $A = O$
  - (3)  $\|\alpha A\| = |\alpha| \|A\|$
  - (4)  $\|A + B\| \leq \|A\| + \|B\|$

- often we add a fifth condition that  $\|\cdot\|$  satisfies the *submultiplicative property*

$$\|AB\| \leq \|A\| \|B\|$$

## 11. HÖLDER NORMS

- example: *Frobenius norm*

$$\|A\|_F = \left( \sum_{i=1}^m \sum_{j=1}^n |a_{ij}|^2 \right)^{1/2}$$

which is submultiplicative since

$$\|AB\|_F^2 = \sum_{i=1}^m \sum_{k=1}^p \left| \sum_{j=1}^n a_{ij} b_{jk} \right|^2 \leq \sum_{i=1}^m \sum_{k=1}^p \left[ \left( \sum_{j=1}^n |a_{ij}|^2 \right) \left( \sum_{j=1}^n |b_{jk}|^2 \right) \right]$$

by the Cauchy-Schwarz inequality and the last expression is equal to

$$\left( \sum_{i=1}^m \sum_{j=1}^n |a_{ij}|^2 \right) \left( \sum_{k=1}^p \sum_{j=1}^n |b_{jk}|^2 \right) = \|A\|_F^2 \|B\|_F^2$$

- example: more generally we have Hölder  $p$ -norm for any  $p \in [1, \infty]$ ,

$$\|A\|_{H,p} = \left( \sum_{i=1}^m \sum_{j=1}^n |a_{ij}|^p \right)^{1/p}$$

and

$$\|A\|_{H,\infty} = \max_{i,j} |a_{ij}|$$

- Hölder norms are obtained by viewing an  $m \times n$  matrix  $A = [a_{ij}]_{i,j=1}^{m,n} \in \mathbb{C}^{m \times n}$  as a vector  $\alpha = [a_{11}, a_{12}, \dots, a_{mn}]^T \in \mathbb{C}^{mn}$  with  $mn$  entries, this is often written as  $\alpha = \text{vec}(A)$
- we have  $\|A\|_{H,p} = \|\text{vec}(A)\|_p$
- clearly  $\|A\|_{H,2} = \|A\|_F = \|\text{vec}(A)\|_2$
- in general Hölder  $p$ -norms are not submultiplicative for  $p \neq 2$ 
  - example: take  $A = \begin{bmatrix} 1 & 1 \\ 0 & 0 \end{bmatrix}$ ,  $B = \begin{bmatrix} 1 & 0 \\ 1 & 0 \end{bmatrix}$ , then  $AB = \begin{bmatrix} 2 & 0 \\ 0 & 0 \end{bmatrix}$  but

$$\|AB\|_{H,\infty} = 2 > 1 = \|A\|_{H,\infty} \|B\|_{H,\infty}$$